

# Monte Carlo Tree Search



Which root node should we select to run simulations on?

### Idea #1: Epsilon Greedy


Select random node with epsilon probability  
Select best node with (1-epsilon) probability

Problem:

If we are very confident a node is bad, then we want 0 probability of expanding it.

Epsilon-greedy treats every action other than the best action equally.

Wins/Total (for player 1)



### Idea #2: Thompson Sampling

Sample action with probability proportional to value of each action (wins/total)

$$p(a_t) \propto e^{w/N}$$

Advantage: Prioritizes sampling actions based on quality (i.e., higher probability of simulating second best action than worst action)

Disadvantage: Does not take into account uncertainty

# Multi-Arm Bandits



# Multi-Arm Bandits

Single-armed bandit



# Multi-Arm Bandits



Single-armed bandit



# Multi-Arm Bandits



Single-armed bandit



# Multi-Arm Bandits



Single-armed bandit



# Multi-Arm Bandits



Single-armed bandit



# Multi-Arm Bandits



Single-armed bandit



# Multi-Arm Bandits



Single-armed bandit



# Multi-Arm Bandits



Single-armed bandit



# Multi-Arm Bandits

Single-armed bandit



# Multi-Arm Bandits

When an arm is pulled, the rewards are random.

Each arm returns a reward with (different) unknown mean and variance

Single-armed bandit



# Multi-Arm Bandits

When an arm is pulled, the rewards are random.

Each arm returns a reward with (different) unknown mean and variance

Bandit Problems are essentially MDPs with a single state

Single-armed bandit



How can get as much reward as possible over N pulls?



# Multi-Arm Bandits

When an arm is pulled, the rewards are random.

Each arm returns a reward with (different) unknown mean and variance

Bandit Problems are essentially MDPs with a single state

Single-armed bandit




Which node to select in MCTS is a Bandit problem!

Each action returns a random result and we'd like to select the best action as frequently as possible.



# Performance of Different Exploration Policies in Multi-Armed Bandits



# Upper Confidence Bound (UCB)

Select action according to:

$$UCB(t) = \operatorname{argmax}_i \hat{\mu}_i + \sqrt{\frac{C \ln(t)}{n_i}}$$

Some constant times the log of number of rounds

Action at round t

Action i that maximizes...

Empirical mean of that action so far

Number of times action i has been executed

# Upper Confidence Bound (UCB)

Select action according to:

$$UCB(t) = \operatorname{argmax}_i \hat{\mu}_i + \sqrt{\frac{C \ln(t)}{n_i}}$$

How good that action is


How “explored” that action is compared to other actions

UCB is near-optimal for Multi-Armed Bandits problems

When  $t$  is large and  $n_i$  is small, this term is larger. Action is more likely to be selected.

Which root node should we select to run simulations on?

Using C=2 (common choice)



$$\text{UCB score} = \frac{7}{10} + \sqrt{\frac{2 \cdot \ln(21)}{10}}$$
$$= \frac{7}{10} + .61$$


$$\text{UCB score} = \frac{3}{8} + \sqrt{\frac{2 \cdot \ln(21)}{8}}$$
$$= \frac{3}{8} + .87$$

$$\text{UCB score} = \frac{0}{3} + \sqrt{\frac{2 \cdot \ln(21)}{3}}$$
$$= 0 + 1.42$$

UCB selects action 2

MCTS will iteratively grow the search tree, but not every branch will be grown at the same rate.

How can we select between root nodes to run simulations?




**Selection**

**Expansion**


Repeated X times

**Simulation**


**Backpropagation**




The selection strategy is applied recursively until an unknown position is reached



One node is added to the tree  
(or more nodes)



One simulated game is played  
(or more games)




The result of this game is backpropagated in the tree

## Selection

At each node in search tree, what action should we take?

Selection strategy balances learning more about best action and learning more about uncertain actions



The selection strategy is applied recursively until an unknown position is reached


# MCTS Tree Nodes

## → Backpropagation

Each node tracks:

1. Number of simulated wins from node or children of node
2. Total number of simulations
3. Parents and Children of node

Backpropagation updates total number of wins and simulations for parent nodes



The result of this game is backpropagated in the tree

---

**Algorithm 1** Monte-Carlo-Tree-Search(state,  $\pi_S$ ) → action

---

$tree \leftarrow MCTSNode(state)$

**while** time-remaining **do**

$leaf \leftarrow SELECT(tree, \pi_S)$

$children \leftarrow EXPAND(leaf)$

$result \leftarrow SIMULATE(children)$

$BACKPROPAGATE(results, children)$

**end while**

Return: Best Action

---

**Algorithm 2** SELECT(state,  $\pi_S$ )

---

*currNode*  $\leftarrow$  state  
**while** *isLeaf*(*currNode*) **do**  
    *currNode*  $\leftarrow \pi_S(\textit{currNode}.\textit{children})$   
    *currNode.visits*  $\leftarrow \textit{currNode}.\textit{visits} + 1$   
**end while**  
**return** *currNode*

---

---

**Algorithm 3** Expand(leaf)

```
children ← []
state ← leaf.state
actions ← state.legalActions()
for action in actions do
    children.append(transition(state, action))
end for
return children
```

---

---

**Algorithm 4** SIMULATE(*children*)

---

```
results ← []
for child in children do
    result = rollout(child)
    results.append(result)
end for
return results
```

---

---

**Algorithm 5** BACKPROPAGATE (results, children)

**Input:** A new leaf node (children) and simulation result for each new leaf node (results)

**for** child in children, result in results **do**

    currNode  $\leftarrow$  child

**while** currNode  $\neq$  NULL **do**

        currNode.visits  $\leftarrow$  currNode.visits +1

**if** result == WHITE-WIN & currNode.player == BLACK **then**

            currNode.value  $\leftarrow$  currNode.value +1

**else if** result == BLACK-WIN & currNode.player == WHITE **then**

            currNode.value  $\leftarrow$  currNode.value +1

**end if**

        currNode  $\leftarrow$  currNode.parent

**end while**

**end for**

---

---

**Algorithm 1** Monte-Carlo-Tree-Search(state,  $\pi_S$ ) → action

---

$tree \leftarrow MCTSNode(state)$

**while** time-remaining **do**

$leaf \leftarrow SELECT(tree, \pi_S)$

$children \leftarrow EXPAND(leaf)$

$result \leftarrow SIMULATE(children)$

$BACKPROPAGATE(results, children)$

**end while**

**return** The action of the node in  $children(tree)$  with the highest number of visits


---

# MCTS Search Tree

(Ideally)

MCTS automatically  
searches branches with  
better expected outcomes


MCTS trees go deeper in  
promising directions and  
remain shallow in poor  
quality branches



# Further Reading


[Monte-Carlo Tree Search](#), Guillaume Chaslot's Dissertation, 2006

# AlphaGo



Learn a Value function and a Policy function, use in MCTS

# MCTS in AlphaGo



# Selection

$$a_t = \underset{a}{\operatorname{argmax}}(Q(s_t, a) + u(s_t, a))$$

$$u(s, a) \propto \frac{P(s, a)}{1 + N(s, a)}$$

Policy Network

$$N(s, a) = \sum_{i=1}^n 1(s, a, i)$$

$$Q(s, a) = \frac{1}{N(s, a)} \sum_{i=1}^n 1(s, a, i) V(s_L^i)$$


Total number of simulations/visits

Average value from  
that state (including  
value estimated with  
Value network and  
result of simulations)

$$V(s_L) = (1 - \lambda)v_\theta(s_L) + \lambda z_L$$

Value Network

Simulation result




# Evaluation


Instead of running random simulations, use a rollout policy  $p_\pi$ .

This is a **simple and fast** rollout policy.

It doesn't need to be perfect, it just needs to be fast. Anything better than random will help.



# Training Value and Policy Networks



# Self-Play

How do you set up a RL environment for Go? Who is your opponent?

→ Yourself!

AlphaGo repeatedly played against itself and improved its parameters using Reinforcement Learning

# AlphaGo

Could defeat the European Champion!

Game 1  
Fan Hui (Black), AlphaGo (White)  
AlphaGo wins by 2.5 points

