

Monte Carlo Tree Search



Which root node should we select to run simulations on?

Idea #1: Epsilon Greedy


Select random node with epsilon probability
Select best node with (1-epsilon) probability

Problem:

If we are very confident a node is bad, then we want 0 probability of expanding it.

Epsilon-greedy treats every action other than the best action equally.

Wins/Total (for player 1)



Idea #2: Thompson Sampling

Sample action with probability proportional to value of each action (wins/total)

$$p(a_t) \propto e^{w/N}$$

Advantage: Prioritizes sampling actions based on quality (i.e., higher probability of simulating second best action than worst action)

Disadvantage: Does not take into account uncertainty

Multi-Arm Bandits



Multi-Arm Bandits

Single-armed bandit



Multi-Arm Bandits



Single-armed bandit



Multi-Arm Bandits



Single-armed bandit



Multi-Arm Bandits



Single-armed bandit



Multi-Arm Bandits



Single-armed bandit



Multi-Arm Bandits



Single-armed bandit



Multi-Arm Bandits



Single-armed bandit



Multi-Arm Bandits



Single-armed bandit



Multi-Arm Bandits

Single-armed bandit



Multi-Arm Bandits

When an arm is pulled, the rewards are random.

Each arm returns a reward with (different) unknown mean and variance

Single-armed bandit



Multi-Arm Bandits

When an arm is pulled, the rewards are random.

Each arm returns a reward with (different) unknown mean and variance

Bandit Problems are essentially MDPs with a single state

Single-armed bandit



How can get as much reward as possible over N pulls?



Multi-Arm Bandits

When an arm is pulled, the rewards are random.

Each arm returns a reward with (different) unknown mean and variance

Bandit Problems are essentially MDPs with a single state

Single-armed bandit




Which node to select in MCTS is a Bandit problem!

Each action returns a random result and we'd like to select the best action as frequently as possible.



Performance of Different Exploration Policies in Multi-Armed Bandits



Upper Confidence Bound (UCB)

Select action according to:

$$UCB(t) = \operatorname{argmax}_i \hat{\mu}_i + \sqrt{\frac{C \ln(t)}{n_i}}$$

Some constant times the log of number of rounds

Action at round t

Action i that maximizes...

Empirical mean of that action so far

Number of times action i has been executed

Upper Confidence Bound (UCB)

Select action according to:

$$UCB(t) = \operatorname{argmax}_i \hat{\mu}_i + \sqrt{\frac{C \ln(t)}{n_i}}$$

How good that action is


How “explored” that action is compared to other actions

UCB is near-optimal for Multi-Armed Bandits problems

When t is large and n_i is small, this term is larger. Action is more likely to be selected.

Which root node should we select to run simulations on?

Using C=2 (common choice)



$$\text{UCB score} = \frac{7}{10} + \sqrt{\frac{2 \cdot \ln(21)}{10}}$$
$$= \frac{7}{10} + .61$$


$$\text{UCB score} = \frac{3}{8} + \sqrt{\frac{2 \cdot \ln(21)}{8}}$$
$$= \frac{3}{8} + .87$$

$$\text{UCB score} = \frac{0}{3} + \sqrt{\frac{2 \cdot \ln(21)}{3}}$$
$$= 0 + 1.42$$

UCB selects action 2

MCTS will iteratively grow the search tree, but not every branch will be grown at the same rate.

How can we select between root nodes to run simulations?






Selection

Expansion

Repeated X times

Simulation

Backpropagation



The selection strategy is applied recursively until an unknown position is reached

One node is added to the tree
(or more nodes)


One simulated game is played
(or more games)

The result of this game is backpropagated in the tree

Selection

At each node in search tree, what action should we take?

Selection strategy balances learning more about best action and learning more about uncertain actions



The selection strategy is applied recursively until an unknown position is reached


MCTS Tree Nodes

→ Backpropagation

Each node tracks:

1. Number of simulated wins from node or children of node
2. Total number of simulations
3. Parents and Children of node

Backpropagation updates total number of wins and simulations for parent nodes



The result of this game is backpropagated in the tree

Algorithm 1 Monte-Carlo-Tree-Search(state, π_S) → action

$tree \leftarrow MCTSNode(state)$

while time-remaining **do**

$leaf \leftarrow SELECT(tree, \pi_S)$

$children \leftarrow EXPAND(leaf)$

$result \leftarrow SIMULATE(children)$

$BACKPROPAGATE(results, children)$

end while

Return: Best Action

Algorithm 2 SELECT(state, π_S)

currNode \leftarrow state
while *isLeaf*(*currNode*) **do**
 currNode $\leftarrow \pi_S(\textit{currNode}.\textit{children})$
 currNode.visits $\leftarrow \textit{currNode}.\textit{visits} + 1$
end while
return *currNode*

Algorithm 3 Expand(leaf)

```
children ← []
state ← leaf.state
actions ← state.legalActions()
for action in actions do
    children.append(transition(state, action))
end for
return children
```

Algorithm 4 SIMULATE(*children*)

```
results ← []
for child in children do
    result = rollout(child)
    results.append(result)
end for
return results
```

Algorithm 5 BACKPROPAGATE (results, children)

Input: A new leaf node (children) and simulation result for each new leaf node (results)

for child in children, result in results **do**

 currNode \leftarrow child

while currNode \neq NULL **do**

 currNode.visits \leftarrow currNode.visits +1

if result == WHITE-WIN & currNode.player == BLACK **then**

 currNode.value \leftarrow currNode.value +1

else if result == BLACK-WIN & currNode.player == WHITE **then**

 currNode.value \leftarrow currNode.value +1

end if

 currNode \leftarrow currNode.parent

end while

end for

Algorithm 1 Monte-Carlo-Tree-Search(state, π_S) → action

$tree \leftarrow MCTSNode(state)$

while time-remaining **do**

$leaf \leftarrow SELECT(tree, \pi_S)$

$children \leftarrow EXPAND(leaf)$

$result \leftarrow SIMULATE(children)$

$BACKPROPAGATE(results, children)$

end while


return The action of the node in $children(tree)$ with the highest number of visits

MCTS Search Tree

(Ideally)

MCTS automatically
searches branches with
better expected outcomes


MCTS trees go deeper in
promising directions and
remain shallow in poor
quality branches



Further Reading


[Monte-Carlo Tree Search](#), Guillaume Chaslot's Dissertation, 2006

AlphaGo



Learn a Value function and a Policy function, use in MCTS

MCTS in AlphaGo



Selection

$$a_t = \underset{a}{\operatorname{argmax}}(Q(s_t, a) + u(s_t, a))$$

$$u(s, a) \propto \frac{P(s, a)}{1 + N(s, a)}$$

Policy Network

$$N(s, a) = \sum_{i=1}^n 1(s, a, i)$$

$$Q(s, a) = \frac{1}{N(s, a)} \sum_{i=1}^n 1(s, a, i) V(s_L^i)$$


Total number of simulations/visits

Average value from
that state (including
value estimated with
Value network and
result of simulations)

$$V(s_L) = (1 - \lambda)v_\theta(s_L) + \lambda z_L$$

Value Network

Simulation result




Evaluation


Instead of running random simulations, use a rollout policy p_π .

This is a **simple and fast** rollout policy.

It doesn't need to be perfect, it just needs to be fast. Anything better than random will help.



Training Value and Policy Networks



Self-Play

How do you set up a RL environment for Go? Who is your opponent?

AlphaGo

Game 1
Fan Hui (Black), AlphaGo (White)
AlphaGo wins by 2.5 points

