



# Policy Gradient Methods

CSCI 410

Eric Ewing



# Deep-Q Networks



Video Source: <https://www.youtube.com/watch?v=V1eYniJ0Rnk>

# Q-Values to Policy

What do we do after we learn Q? We need to turn them into a policy.

For a given state, take the action associated with the best Q-value.


$$\pi(s) = \operatorname{argmax}_a Q(s, a)$$

# Policies

Why learn Q-values first and turn them into a policy? Why not just learn a policy?


# Policies

Why learn Q-values first and turn them into a policy? Why not just learn a policy?




# Policies

Why learn Q-values first and turn them into a policy? Why not just learn a policy?



# Policies

Why learn Q-values first and turn them into a policy? Why not just learn a policy?



# How do we train a policy network?

# How do we train a policy network?

Need to find an appropriate loss function.

# How do we train a policy network?

Need to find an appropriate loss function.

What's our objective?

# How do we train a policy network?

Need to find an appropriate loss function.

What's our objective?

Find a policy  $\pi$  such that the value of the start state is maximized:

# How do we train a policy network?

Need to find an appropriate loss function.

What's our objective?

Find a policy  $\pi$  such that the value of the start state is maximized:

$$\pi = \operatorname{argmax}_{\pi} (V(s_0))$$

# How do we train a policy network?

Need to find an appropriate loss function.

What's our objective?

Find a policy  $\pi$  such that the value of the start state is maximized:

$$\pi = \operatorname{argmax}_{\pi} (V(s_0))$$

How can we maximize  $V(s_0)$ ?

Let  $J(\theta)$  be our objective function:

$$J(\theta) = V(s_0)$$

Let  $J(\theta)$  be our objective function:

$$J(\theta) = V(s_0)$$

$$J(\theta) = \mathbb{E}[G_0]$$

Let  $J(\theta)$  be our objective function:

$$J(\theta) = V(s_0)$$

$$J(\theta) = \mathbb{E}[G_0]$$

$$J(\theta) = \sum_{\tau} \Pr(\tau|\theta) G(\tau)$$

Let  $J(\theta)$  be our objective function:

$$J(\theta) = V(s_0)$$

$$J(\theta) = \mathbb{E}[G_0]$$

$$J(\theta) = \sum_{\tau} \Pr(\tau|\theta) G(\tau)$$

Probability of a trajectory occurring

Returns of a specific trajectory

Let  $J(\theta)$  be our objective function:

$$J(\theta) = V(s_0)$$

$$J(\theta) = \mathbb{E}[G_0]$$

$$J(\theta) = \sum_{\tau} \Pr(\tau|\theta) G(\tau)$$

↗

Probability of a trajectory occurring

Returns of a specific trajectory

$$\Pr(\tau|\theta) = \prod_{t=0}^T P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)$$

Let  $J(\theta)$  be our objective function:

$$J(\theta) = V(s_0)$$

$$J(\theta) = \mathbb{E}[G_0]$$

$$J(\theta) = \sum_{\tau} \Pr(\tau|\theta) G(\tau)$$

Probability of a trajectory occurring

Returns of a specific trajectory

$$\Pr(\tau|\theta) = \prod_{t=0}^T P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)$$

State transition Probability

Probability of taking an action for a given state

Let  $J(\theta)$  be our objective function:

$$J(\theta) = V(s_0)$$

$$J(\theta) = \mathbb{E}[G_0]$$

$$J(\theta) = \sum_{\tau} \Pr(\tau|\theta) G(\tau)$$

Sum over all possible trajectories

Probability of a trajectory occurring

Returns of a specific trajectory

$$\Pr(\tau|\theta) = \prod_{t=0}^T P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)$$

State transition Probability

Probability of taking an action for a given state

# Log-Derivative Trick

We can rewrite the derivative of a function using the derivative of the natural log function:

$$\nabla \ln f(x) = \frac{\nabla f(x)}{f(x)}$$

$$\nabla f(x) = f(x) \nabla \ln f(x)$$

When applied to  $\Pr(\tau|\theta)$ :

$$\nabla_{\theta} \Pr(\tau|\theta) = \Pr(\tau|\theta) \nabla_{\theta} \ln \Pr(\tau|\theta)$$

# Log Probability Trick

$$\Pr(\tau|\theta) = \prod_{t=0}^T P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)$$

# Log Probability Trick

$$\Pr(\tau|\theta) = \prod_{t=0}^T P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)$$

$$\nabla_\theta \Pr(\tau|\theta) = \Pr(\tau|\theta) \nabla_\theta \ln \Pr(\tau|\theta)$$

This gradient term is  
what we want to  
calculate

# Log Probability Trick

$$\Pr(\tau|\theta) = \prod_{t=0}^T P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)$$

$$\nabla_\theta \Pr(\tau|\theta) = \Pr(\tau|\theta) \nabla_\theta \ln \Pr(\tau|\theta)$$

This gradient term is  
what we want to  
calculate

$$\nabla_\theta \ln \Pr(\tau|\theta) = \nabla_\theta \sum_{t=0}^T \ln[P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)]$$

# Log Probability Trick

$$\Pr(\tau|\theta) = \prod_{t=0}^T P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)$$

$$\nabla_\theta \Pr(\tau|\theta) = \Pr(\tau|\theta) \nabla_\theta \ln \Pr(\tau|\theta)$$

This gradient term is what we want to calculate

$$\nabla_\theta \ln \Pr(\tau|\theta) = \nabla_\theta \sum_{t=0}^T \ln [P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)]$$

$$\nabla_\theta \ln \Pr(\tau|\theta) = \nabla_\theta \sum_{t=0}^T \ln P(s_{t+1}|s_t, a_t) + \ln \pi_\theta(a_t|s_t)$$

# Log Probability Trick

$$\Pr(\tau|\theta) = \prod_{t=0}^T P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)$$

$$\nabla_\theta \Pr(\tau|\theta) = \Pr(\tau|\theta) \nabla_\theta \ln \Pr(\tau|\theta)$$

This gradient term is what we want to calculate

$$\nabla_\theta \ln \Pr(\tau|\theta) = \nabla_\theta \sum_{t=0}^T \ln P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)$$

Log of product -> sum of logs

$$\nabla_\theta \ln \Pr(\tau|\theta) = \nabla_\theta \sum_{t=0}^T \ln P(s_{t+1}|s_t, a_t) + \ln \pi_\theta(a_t|s_t)$$

Log of product -> sum of logs

$$\nabla_\theta \ln \Pr(\tau|\theta) = \sum_{t=0}^T \nabla_\theta \ln P(s_{t+1}|s_t, a_t) + \nabla_\theta \ln \pi_\theta(a_t|s_t)$$

Derivative of sum -> sum of derivative

# Gradient of a trajectory

$$\nabla_{\theta} \ln \Pr(\tau|\theta) = \sum_{t=0}^T \nabla_{\theta} \ln P(s_{t+1}|s_t, a_t) + \nabla_{\theta} \ln \pi_{\theta}(a_t|s_t)$$



State transition function  
does not depend on  $\theta$ !

$$\nabla_{\theta} \ln \Pr(\tau|\theta) = \sum_{t=0}^T \nabla_{\theta} \ln \pi_{\theta}(a_t|s_t)$$

# Policy Gradient Derivation

Putting it all back together:

$$J(\theta) = \sum_{\tau} \Pr(\tau|\theta) G(\tau)$$

Our Objective

# Policy Gradient Derivation

Putting it all back together:

$$J(\theta) = \sum \Pr(\tau|\theta) G(\tau)$$

Our Objective

$$\nabla_{\theta} J(\theta) = \sum_{\tau} \nabla_{\theta} \Pr(\tau|\theta) G(\tau)$$

Take the gradient

# Policy Gradient Derivation

Putting it all back together:

$$J(\theta) = \sum \Pr(\tau|\theta) G(\tau)$$

Our Objective

$$\nabla_{\theta} J(\theta) = \sum_{\tau} \nabla_{\theta} \Pr(\tau|\theta) G(\tau)$$

Take the gradient

$$\nabla_{\theta} J(\theta) = \sum_{\tau} \Pr(\tau|\theta) G(\tau) \nabla_{\theta} \ln \Pr(\tau|\theta)$$

Log-Derivative Trick

# Policy Gradient Derivation

Putting it all back together:

$$J(\theta) = \sum_{\tau} \Pr(\tau|\theta) G(\tau) \quad \text{Our Objective}$$

$$\nabla_{\theta} J(\theta) = \sum_{\tau} \nabla_{\theta} \Pr(\tau|\theta) G(\tau) \quad \text{Take the gradient}$$

$$\nabla_{\theta} J(\theta) = \sum_{\tau} \Pr(\tau|\theta) G(\tau) \nabla_{\theta} \ln \Pr(\tau|\theta) \quad \text{Log-Derivative Trick}$$

$$\nabla_{\theta} J(\theta) = \sum_{\tau} [\Pr(\tau|\theta) G(\tau) \sum_{t=0}^T \nabla_{\theta} \ln \pi_{\theta}(a_t|s_t)] \quad \text{Gradient of a Trajectory}$$

# Policy Gradient Derivation

Putting it all back together:

$$J(\theta) = \sum_{\tau} \Pr(\tau|\theta) G(\tau) \quad \text{Our Objective}$$

$$\nabla_{\theta} J(\theta) = \sum_{\tau} \nabla_{\theta} \Pr(\tau|\theta) G(\tau) \quad \text{Take the gradient}$$

$$\nabla_{\theta} J(\theta) = \sum_{\tau} \Pr(\tau|\theta) G(\tau) \nabla_{\theta} \ln \Pr(\tau|\theta) \quad \text{Log-Derivative Trick}$$

$$\nabla_{\theta} J(\theta) = \sum_{\tau} [\Pr(\tau|\theta) G(\tau) \sum_{t=0}^T \nabla_{\theta} \ln \pi_{\theta}(a_t|s_t)] \quad \text{Gradient of a Trajectory}$$

$$\nabla_{\theta} J(\theta) = \mathbb{E}[G_0 \sum_{t=0}^T \nabla_{\theta} \ln \pi_{\theta}(a_t|s_t)] \quad \text{Convert back to Expectation}$$

# Policy Gradient

Bigger step if better returns

Direction to move in to increase probability of trajectory

$$\nabla_{\theta} J(\theta) = \mathbb{E}[G_0 \sum_{t=0}^T \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t)]$$

We will never be able to sum over all possible trajectories...

How do we get around this?

# Policy Gradient

Bigger step if better returns

Direction to move in to increase probability of trajectory

$$\nabla_{\theta} J(\theta) = \mathbb{E}[G_0 \sum_{t=0}^T \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t)]$$

We will never be able to sum over all possible trajectories...

How do we get around this?

Sampling!

1. Collect n trajectories following policy  $\pi_{\theta}$
2.  $\Pr(\tau|\theta) = 1/n$  for each trajectory
3. Calculate the total return for each trajectory  $G(\tau)$

# Reward-To-Go Policy Gradient

You can also do the policy gradient derivation such that the gradient does not depend on  $G_0$ , but on  $G_t$

$$\nabla_{\theta} J(\theta) = \mathbb{E} \left[ \sum_{t=0}^T G_t \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) \right]$$

Or

$$\nabla_{\theta} J(\theta) = \mathbb{E} \left[ \sum_{t=0}^T Q(s_t, a_t) \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) \right]$$

# REINFORCE (Policy Gradient Learning)

REINFORCE, A Monte-Carlo Policy-Gradient Method (episodic)

Input: a differentiable policy parameterization  $\pi(a|s, \theta)$

Initialize policy parameter  $\theta \in \mathbb{R}^{d'}$

Repeat forever:

    Generate an episode  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$ , following  $\pi(\cdot| \cdot, \theta)$


    For each step of the episode  $t = 0, \dots, T - 1$ :

$G \leftarrow$  return from step  $t$

$\theta \leftarrow \theta + \alpha \gamma^t G \nabla_{\theta} \ln \pi(A_t | S_t, \theta)$

# REINFORCE Variance

If we could calculate  $\nabla_{\theta}J(\theta)$  exactly (not just for single trajectory/sample), then Policy Gradient would be a great algorithm! (with some minor flaws)



Results on Cartpole

# Actor-Critic Methods

# Actor-Critic Methods

REINFORCE uses the return for a trajectory  $G_t$ :

# Actor-Critic Methods

REINFORCE uses the return for a trajectory  $G_t$ :

$$\nabla_{\theta} J(\theta) = \mathbb{E}\left[\sum_{t=0}^T G_t \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t)\right]$$

# Actor-Critic Methods

REINFORCE uses the return for a trajectory  $G_t$ :

$$\nabla_{\theta} J(\theta) = \mathbb{E}\left[\sum_{t=0}^T G_t \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t)\right]$$

Variance of Returns is  
always a problem...

# Actor-Critic Methods

REINFORCE uses the return for a trajectory  $G_t$ :

$$\nabla_{\theta} J(\theta) = \mathbb{E}\left[\sum_{t=0}^T G_t \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t)\right]$$

Variance of Returns is  
always a problem...

Actor-Critic Methods learn an approximation of  $G_t$

# Actor-Critic Methods

REINFORCE uses the return for a trajectory  $G_t$ :

$$\nabla_{\theta} J(\theta) = \mathbb{E}\left[\sum_{t=0}^T G_t \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t)\right]$$

Variance of Returns is  
always a problem...

Actor-Critic Methods learn an approximation of  $G_t$

$$Q^{\pi}(s_t, a_t) = \mathbb{E}[G_t]$$

# Actor-Critic Methods

REINFORCE uses the return for a trajectory  $G_t$ :

$$\nabla_{\theta} J(\theta) = \mathbb{E}\left[\sum_{t=0}^T G_t \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t)\right]$$

Variance of Returns is  
always a problem...

Actor-Critic Methods learn an approximation of  $G_t$

$$Q^{\pi}(s_t, a_t) = \mathbb{E}[G_t]$$

# Actor-Critic Methods

REINFORCE uses the return for a trajectory  $G_t$ :

$$\nabla_{\theta} J(\theta) = \mathbb{E}\left[\sum_{t=0}^T G_t \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t)\right]$$

Variance of Returns is  
always a problem...

Actor-Critic Methods learn an approximation of  $G_t$

$$Q^{\pi}(s_t, a_t) = \mathbb{E}[G_t]$$

$$\nabla_{\theta} J(\theta) = \mathbb{E}\left[\sum_{t=0}^T Q^{\pi_{\theta}}(s_t, a_t) \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t)\right]$$

# Actor-Critic Algorithm: Learn $Q^\pi$ and $\pi(a|s)$

# Actor-Critic Algorithm: Learn $Q^\pi$ and $\pi(a|s)$

Actor (policy): Takes actions



# Actor-Critic Algorithm: Learn $Q^\pi$ and $\pi(a|s)$

Actor (policy): Takes actions



Critic: Scores the action



# Actor-Critic Algorithm: Learn $Q^\pi$ and $\pi(a|s)$

Actor (policy): Takes actions



Critic: Scores the action



# Actor-Critic Algorithm: Learn $Q^\pi$ and $\pi(a|s)$

Initialize  $\pi_\theta, Q_w, \alpha_\theta, \alpha_w$

Policy network has parameters  $\theta$   
Q network has parameters  $w$

Repeat forever:

Take action  $a$ , get new state  $s'$  and reward  $r$

Sample next action  $a' \sim \pi_\theta(a|s)$

update  $\theta \leftarrow \theta + \alpha_\theta Q_w(s, a) \nabla_\theta \ln \pi_\theta(a|s)$

Calculate TD Error:  $\delta = r + \gamma Q_w(s', a') - Q_w(s, a)$

update  $w \leftarrow w + \alpha_w \delta \nabla_w Q_w(s, a)$

$a \leftarrow a', s \leftarrow s'$

# Actor-Critic Algorithm: Learn $Q^\pi$ and $\pi(a|s)$

Initialize  $\pi_\theta, Q_w, \alpha_\theta, \alpha_w$

Policy network has parameters  $\theta$   
Q network has parameters  $w$

Repeat forever:

    Take action  $a$ , get new state  $s'$  and reward  $r$

    Sample next action  $a' \sim \pi_\theta(a|s)$

    update  $\theta \leftarrow \theta + \alpha_\theta Q_w(s, a) \nabla_\theta \ln \pi_\theta(a|s)$

    Calculate TD Error:  $\delta = r + \gamma Q_w(s', a') - Q_w(s, a)$

    update  $w \leftarrow w + \alpha_w \delta \nabla_w Q_w(s, a)$

$a \leftarrow a', s \leftarrow s'$

Like Q-learning and REINFORCE at the same time


# Variations on a Theme...

How to estimate  $J(\theta)$

(Wikipedia uses  $R_t$  instead of  $G_t$ )

- $\sum_{0 \leq i \leq T} (\gamma^i R_i)$ .
- $\gamma^j \sum_{j \leq i \leq T} (\gamma^{i-j} R_i)$ : the **REINFORCE** algorithm.
- $\gamma^j \sum_{j \leq i \leq T} (\gamma^{i-j} R_i) - b(S_j)$ : the **REINFORCE with baseline** algorithm. Here  $b$  is an arbitrary function.
- $\gamma^j (R_j + \gamma V^{\pi_\theta}(S_{j+1}) - V^{\pi_\theta}(S_j))$ : **TD(1)** learning.
- $\gamma^j Q^{\pi_\theta}(S_j, A_j)$ .
- $\gamma^j A^{\pi_\theta}(S_j, A_j)$ : **Advantage Actor-Critic (A2C)**.<sup>[3]</sup>
- $\gamma^j (R_j + \gamma R_{j+1} + \gamma^2 V^{\pi_\theta}(S_{j+2}) - V^{\pi_\theta}(S_j))$ : TD(2) learning.
- $\gamma^j \left( \sum_{k=0}^{n-1} \gamma^k R_{j+k} + \gamma^n V^{\pi_\theta}(S_{j+n}) - V^{\pi_\theta}(S_j) \right)$ : TD(n) learning.
- $\gamma^j \sum_{n=1}^{\infty} \frac{\lambda^{n-1}}{1-\lambda} \cdot \left( \sum_{k=0}^{n-1} \gamma^k R_{j+k} + \gamma^n V^{\pi_\theta}(S_{j+n}) - V^{\pi_\theta}(S_j) \right)$ : TD( $\lambda$ ) learning, also known as **GAE (generalized advantage estimate)**.<sup>[4]</sup> This is obtained by an exponentially decaying sum of the TD(n) learning terms.

# Actor-Critic Networks



# Actor-Critic Networks



Just make sure you use the correct activation function for the different outputs

# Deep Q-Learning Revisited

Compute TD-Error:  $\delta = r + \gamma \max_{a'} Q(s', a') - Q(s, a)$

Loss Function:  $L = \delta^2$

Update model with Gradient Descent

# Deep Q-Learning Revisited

Compute TD-Error:  $\delta = r + \gamma \max_{a'} Q(s', a') - Q(s, a)$

Loss Function:  $L = \delta^2$

Update model with Gradient Descent

Q-Learning uses:

$$\delta = r + \gamma \max_{a'} Q(s', a') - Q(s, a)$$

Actor-Critic Uses:

$$\delta = r + \gamma Q(s', a') - Q(s, a)$$

# Deep Q-Learning Revisited

Compute TD-Error:  $\delta = r + \gamma \max_{a'} Q(s', a') - Q(s, a)$

Loss Function:  $L = \delta^2$

Update model with Gradient Descent

Q-Learning uses:

$$\delta = r + \gamma \max_{a'} Q(s', a') - Q(s, a)$$

Actor-Critic Uses:

$$\delta = r + \gamma Q(s', a') - Q(s, a)$$

Q-Learning is learning *Optimal Q-values*

Actor-Critic is learning the Q-values for  
following a specific policy  $Q^\pi$

# On-Policy Vs. Off-Policy Learning

RL algorithms collect experiences and learn from these experiences

*On-Policy Algorithms* have to collect experiences with the policy they are learning

*Off-Policy Algorithms* can use **any** policy to collect experiences

# DQNs Are Off-Policy

In Q-Learning, we typically collect experiences using an  $\epsilon$ -greedy policy in training

With probability  $\epsilon$ , take random action.

Else, take action  $\text{argmax}_a Q(s, a)$

At test time, we should take the best actions, not the  $\epsilon$ -greedy actions

$\text{argmax}_a Q(s, a)$

# DQNs Are Off-Policy

In Q-Learning, we typically collect experiences using an  $\epsilon$ -greedy policy in training

With probability  $\epsilon$ , take random action.

Else, take action  $\text{argmax}_a Q(s, a)$

At test time, we should take the best actions, not the  $\epsilon$ -greedy actions

$\text{argmax}_a Q(s, a)$

These are different policies! DQNs can be trained with any data collection policy at training time

# Actor-Critic Algorithm: Learn $Q^\pi$ and $\pi(a|s)$

Initialize  $\pi_\theta, Q_w, \alpha_\theta, \alpha_w$

Repeat forever:

    Take action  $a$ , get new state  $s'$  and reward  $r$

    Sample next action  $a' \sim \pi_\theta(a|s)$

        On Policy: Have to take actions according to  $\pi_\theta$

    update  $\theta \leftarrow \theta + \alpha_\theta Q_w(s, a) \nabla_\theta \ln \pi_\theta(a|s)$

    Calculate TD Error:  $\delta = r + \gamma Q_w(s', a') - Q_w(s, a)$

    update  $w \leftarrow w + \alpha_w \delta \nabla_w Q_w(s, a)$

$a \leftarrow a', s \leftarrow s'$

# On-Policy vs Off-Policy Learning

# On-Policy vs Off-Policy Learning

**Advantage** of On-Policy Learning:

# On-Policy vs Off-Policy Learning

**Advantage** of On-Policy Learning:

- More sample efficient, tend to converge faster

# On-Policy vs Off-Policy Learning

**Advantage** of On-Policy Learning:

- More sample efficient, tend to converge faster

**Disadvantages** of On-Policy Learning:

# On-Policy vs Off-Policy Learning

## **Advantage** of On-Policy Learning:

- More sample efficient, tend to converge faster

## **Disadvantages** of On-Policy Learning:

- Can get stuck in local minima
  - Is there a simple policy that performs ok? Would small changes to that policy cause returns to go down temporarily?
  - How do balance exploration in our policy?

# On-Policy vs Off-Policy Learning

**Advantage** of On-Policy Learning:

- More sample efficient, tend to converge faster

**Disadvantages** of On-Policy Learning:

- Can get stuck in local minima
  - Is there a simple policy that performs ok? Would small changes to that policy cause returns to go down temporarily?
  - How do balance exploration in our policy?

**Advantages** of Off-Policy Learning:

# On-Policy vs Off-Policy Learning

## **Advantage** of On-Policy Learning:

- More sample efficient, tend to converge faster

## **Disadvantages** of On-Policy Learning:

- Can get stuck in local minima
  - Is there a simple policy that performs ok? Would small changes to that policy cause returns to go down temporarily?
  - How do balance exploration in our policy?

## **Advantages** of Off-Policy Learning:

- Can learn from any policy

# On-Policy vs Off-Policy Learning

## **Advantage** of On-Policy Learning:

- More sample efficient, tend to converge faster

## **Disadvantages** of On-Policy Learning:

- Can get stuck in local minima
  - Is there a simple policy that performs ok? Would small changes to that policy cause returns to go down temporarily?
  - How do balance exploration in our policy?

## **Advantages** of Off-Policy Learning:

- Can learn from any policy
- More likely to learn an optimal policy

# On-Policy vs Off-Policy Learning

## **Advantage** of On-Policy Learning:

- More sample efficient, tend to converge faster

## **Disadvantages** of On-Policy Learning:

- Can get stuck in local minima
  - Is there a simple policy that performs ok? Would small changes to that policy cause returns to go down temporarily?
  - How do balance exploration in our policy?

## **Advantages** of Off-Policy Learning:

- Can learn from any policy
- More likely to learn an optimal policy

## **Disadvantages** of Off-Policy Learning:

# On-Policy vs Off-Policy Learning

## **Advantage** of On-Policy Learning:

- More sample efficient, tend to converge faster

## **Disadvantages** of On-Policy Learning:

- Can get stuck in local minima
  - Is there a simple policy that performs ok? Would small changes to that policy cause returns to go down temporarily?
  - How do balance exploration in our policy?

## **Advantages** of Off-Policy Learning:

- Can learn from any policy
- More likely to learn an optimal policy

## **Disadvantages** of Off-Policy Learning:

- Slower...

# Off-Policy Learning

Most of the time in RL, collecting the data is computationally expensive.

So far, we've looked at an example, learned from it, and discarded it.

In all our other problems, we always learned from data multiple times (i.e., epochs)

# Off-Policy Learning

Most of the time in RL, collecting the data is computationally expensive.

So far, we've looked at an example, learned from it, and discarded it.

In all our other problems, we always learned from data multiple times (i.e., epochs)

# Off-Policy Learning

Most of the time in RL, collecting the data is computationally expensive.

Maybe we shouldn't throw away useful data immediately...

So far, we've looked at an example, learned from it, and discarded it.


In all our other problems, we always learned from data multiple times (i.e., epochs)

# Experience Replay and Replay Buffers


Keep a memory of experiences  
(state, action, reward,  
next\_state)

As you collect new  
experiences, remove oldest  
experiences from buffer

To train model, sample batch  
of data from buffer




# On-Policy Learning



# On-Policy Learning


Can we use Replay Buffers with On-Policy learning algorithms (e.g., REINFORCE, Actor-Critic, etc.)?



# On-Policy Learning

Can we use Replay Buffers with On-Policy learning algorithms (e.g., REINFORCE, Actor-Critic, etc.)?


No! Data in the buffer was collected with an older policy and we can only learn on experiences collected using the current policy...



# On-Policy Learning

Can we use Replay Buffers with On-Policy learning algorithms (e.g., REINFORCE, Actor-Critic, etc.)?

No! Data in the buffer was collected with an older policy and we can only learn on experiences collected using the current policy...



# But what if we actually could...

Off-Policy Policy Gradient:

Data collected under policy  $\beta(a|s)$  (i.e., older version of policy)

We can re-weight our gradient according to the old policy:

$$\rho = \frac{\pi(a|s)}{\beta(a|s)}$$
$$\nabla_{\theta} J(\theta) = \sum_{(s,a) \in batch} \rho \cdot Q^{\pi}(s, a) \nabla_{\theta} \ln \pi(s, a)$$

Actor-Critic with Importance Sampling

# But what if we actually could...

Off-Policy Policy Gradient:

Data collected under policy  $\beta(a|s)$  (i.e., older version of policy)

We can re-weight our gradient according to the old policy:

$$\nabla_{\theta} J(\theta) = \sum_{(s,a) \in batch} \rho \cdot Q^{\pi}(s, a) \nabla_{\theta} \ln \pi(s, a)$$
$$\rho = \frac{\pi(a|s)}{\beta(a|s)}$$

Actor-Critic with Importance Sampling

Store action probabilities  $\beta(a|s)$  in replay buffer

# Trust Region Policy Optimization

Insight: the reason that variance is bad is that it can cause large updates to  $\pi_\theta$

Add a constraint to how large of an update can be applied:

KL-Divergence between old and new policy must be below some hyperparameter  $\Delta$

$$D_{KL}(\pi_\theta^{new}(\cdot | s) \mid\mid \pi_\theta^{old}(\cdot | s)) \leq \Delta$$

$$\rho = \frac{\pi^{new}(a|s)}{\pi^{old}(a|s)}$$

$$J^{TRPO}(\theta) = \mathbb{E} \left[ \rho \cdot (r + \gamma V^{\pi^{old}}(s') - V^{\pi^{old}}(s)) \right]$$

# Proximal Policy Optimization

TRPO is complicated...

What if instead of constraining the update with KL-Divergence, we clipped the update if it's too big...


$$\rho_{clipped} = clip\left[\frac{\pi^{new}(a|s)}{\pi^{old}(a|s)}, 1 - \epsilon, 1 + \epsilon\right]$$

$$J^{PPO}(\theta) = \mathbb{E}[\min(\rho_{clipped} \cdot \left(r + \gamma V^{\pi^{old}}(s') - V^{\pi^{old}}(s)\right), \rho \left(r + \gamma V^{\pi^{old}}(s') - V^{\pi^{old}}(s)\right))]$$

# PPO

PPO is (basically) State-Of-The-Art (SOTA)


Provides **fast**, **sample-efficient**, and **stable** training



# PPO

PPO is (basically) State-Of-The-Art (SOTA)

Provides **fast**, **sample-efficient**, and **stable** training



# PPO: OpenAI5



# PPO

## Final phase of training ChatGPT

Step 3

Optimize a policy against the reward model using the PPO reinforcement learning algorithm.


A new prompt is sampled from the dataset.

The PPO model is initialized from the supervised policy.

The policy generates an output.

The reward model calculates a reward for the output.

The reward is used to update the policy using PPO.



# Group Relative Policy Optimization

PPO still needs a value estimate  $V^\pi$ , what if we had another way of estimating  $V^\pi$

More specifically, for each question  $q$ , GRPO samples a group of outputs  $\{o_1, o_2, \dots, o_G\}$  from the old policy  $\pi_{\theta_{old}}$  and then optimizes the policy model by maximizing the following objective:

$$\begin{aligned} \mathcal{J}_{GRPO}(\theta) = & \mathbb{E}[q \sim P(Q), \{o_i\}_{i=1}^G \sim \pi_{\theta_{old}}(O|q)] \\ & \frac{1}{G} \sum_{i=1}^G \frac{1}{|o_i|} \sum_{t=1}^{|o_i|} \left\{ \min \left[ \frac{\pi_\theta(o_{i,t}|q, o_{i,<t})}{\pi_{\theta_{old}}(o_{i,t}|q, o_{i,<t})} \hat{A}_{i,t}, \text{clip} \left( \frac{\pi_\theta(o_{i,t}|q, o_{i,<t})}{\pi_{\theta_{old}}(o_{i,t}|q, o_{i,<t})}, 1 - \varepsilon, 1 + \varepsilon \right) \hat{A}_{i,t} \right] - \beta \mathbb{D}_{KL} [\pi_\theta || \pi_{ref}] \right\}, \end{aligned} \quad (3)$$

where  $\varepsilon$  and  $\beta$  are hyper-parameters, and  $\hat{A}_{i,t}$  is the advantage calculated based on relative rewards of the outputs inside each group only.